

Linear Regression

Wednesday, April 5, 2017 9:16 PM

Linear regression . fit

$$\hat{\beta} = \underbrace{(X^T X)^{-1}}_{\substack{1. \text{ matrix multiply} \\ 2. \text{ matrix solve}}} X^T y \quad \left. \vphantom{\hat{\beta}} \right\} \begin{array}{l} \text{matrix} \\ \text{multiply} \end{array}$$

1. Matrix multiply (X is $n \times p$)

$$(X^T X)_{jk} = \sum_i X_{ij} X_{ik} \quad - \quad O(n)$$

Computational complexity :

$$O(p^2 n)$$

2. Matrix solve - Cholesky decomp.

Computational complexity:

$$O(p^3)$$

Total Regression Fit Complexity :

$$O(p^3 + p^2 n)$$

Linear Regression, predict

$$\hat{f}(x^*) = x^{*T} \hat{\beta}$$

$O(p)$ time

▷ predict is fast $O(p)$

fit is slow $O(p^3 + p^2n)$

▷ $X^T X$ needs to be invertible
(impossible if $p > n$!)

▷ \hat{f} is linear

▷ Gauss - Markov theorem

Let $y_i = x_i^T \beta + \varepsilon_i$ with $E\varepsilon_i = 0$


$V\varepsilon_i = \sigma^2$ then OLS is

unbiased: $E\hat{\beta} = \beta$

and is the unbiased estimator
with minimum variance.

K Nearest Neighbors

Wednesday, April 5, 2017 10:14 PM



Given a metric $d(x, x')$ then
the k -nearest neighbors of x in
 $\{x_i\}_{i=1}^n$ is x_{m_1}, \dots, x_{m_k} s.t.

$$d(x_{m_1}, x) \leq d(x_{m_2}, x) \leq \dots \leq d(x_{m_n}, x)$$

kNN methods fit a model for
predicting x using the kNN
dataset $\{x_{m_j}, y_{m_j}\}_{j=1}^k$.


Regression: $\hat{f}(x) = \frac{1}{k} \sum_{j=1}^k y_{m_j}$

Classification: $\hat{f}(x) = 1\left\{\frac{1}{k} \sum_{j=1}^k y_{m_j} > \frac{1}{2}\right\}$

▷ randomized rounding if $= \frac{1}{2}$

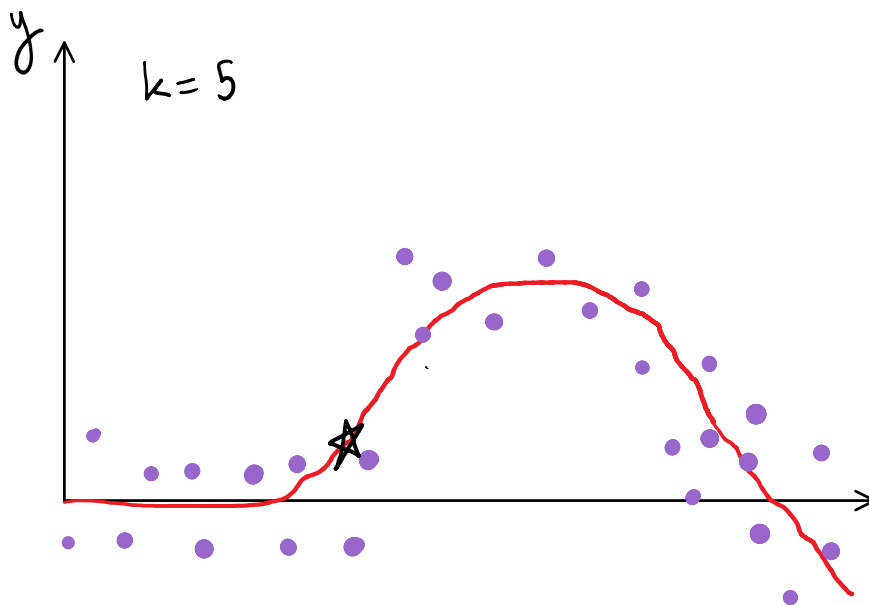
y_{m_1}, \dots, y_{m_k}

distance from $y = -2$



low bias

high variance



high bias

low variance

$$\{y_i, x_i\} \text{ i.i.d}$$

$$\mathbb{E} \left[\frac{1}{k} \sum_{j=1}^k y_{m_j} | x \right] = \frac{1}{k} \sum_{j=1}^k \mathbb{E}[y | x_{m_j}]$$

$$\mathbb{V} \left[\frac{1}{k} \sum_{j=1}^k y_{m_j} | x \right] = \frac{1}{k^2} \sum_{j=1}^k \mathbb{V}[y | x_{m_j}] = \frac{1}{k} \sigma^2$$

$$\text{Bias: } \mathbb{E}[\hat{f}(x)] - \mathbb{E}[y|x] = \frac{1}{k} \sum_{j=1}^k \mathbb{E}[y | x_{m_j}] - \mathbb{E}[y|x]$$

$$\text{as } d(x_{m_j}, x) \uparrow \quad |\mathbb{E}[y | x_{m_j}] - \mathbb{E}[y|x]| \uparrow$$

So, typically as $k \uparrow$ bias \uparrow
and

as $k \uparrow$ Variance \downarrow

this is the bias - variance trade off


$$\text{MSE} = \mathbb{E}(\hat{f}(x) - \mathbb{E}[y|x])^2 \quad (\text{True Risk})$$

$$= \mathbb{E}(\hat{f}(x) - \mathbb{E}\hat{f}(x) + \mathbb{E}\hat{f}(x) - \mathbb{E}[y|x])^2$$

$$= \mathbb{E}(\hat{f}(x) - \mathbb{E}\hat{f}(x))^2 + (\mathbb{E}\hat{f}(x) - \mathbb{E}[y|x])^2$$


$$+ 2 \underbrace{\mathbb{E}(\hat{f}(x) - \mathbb{E}\hat{f}(x))}_0 \cdot (\mathbb{E}\hat{f}(x) - \mathbb{E}[y|x])$$

$$= \text{bias}^2 + \text{Var} \hat{f}(x)$$



ESL Ch. 2

Linear Regression of 0/1 Response






FIGURE 2.1. A classification example in two dimensions. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1), and then fit by linear regression. The line is the decision boundary defined by $x^T \hat{\beta} = 0.5$. The orange shaded region denotes that part of input space classified as ORANGE, while the blue region is classified as BLUE.

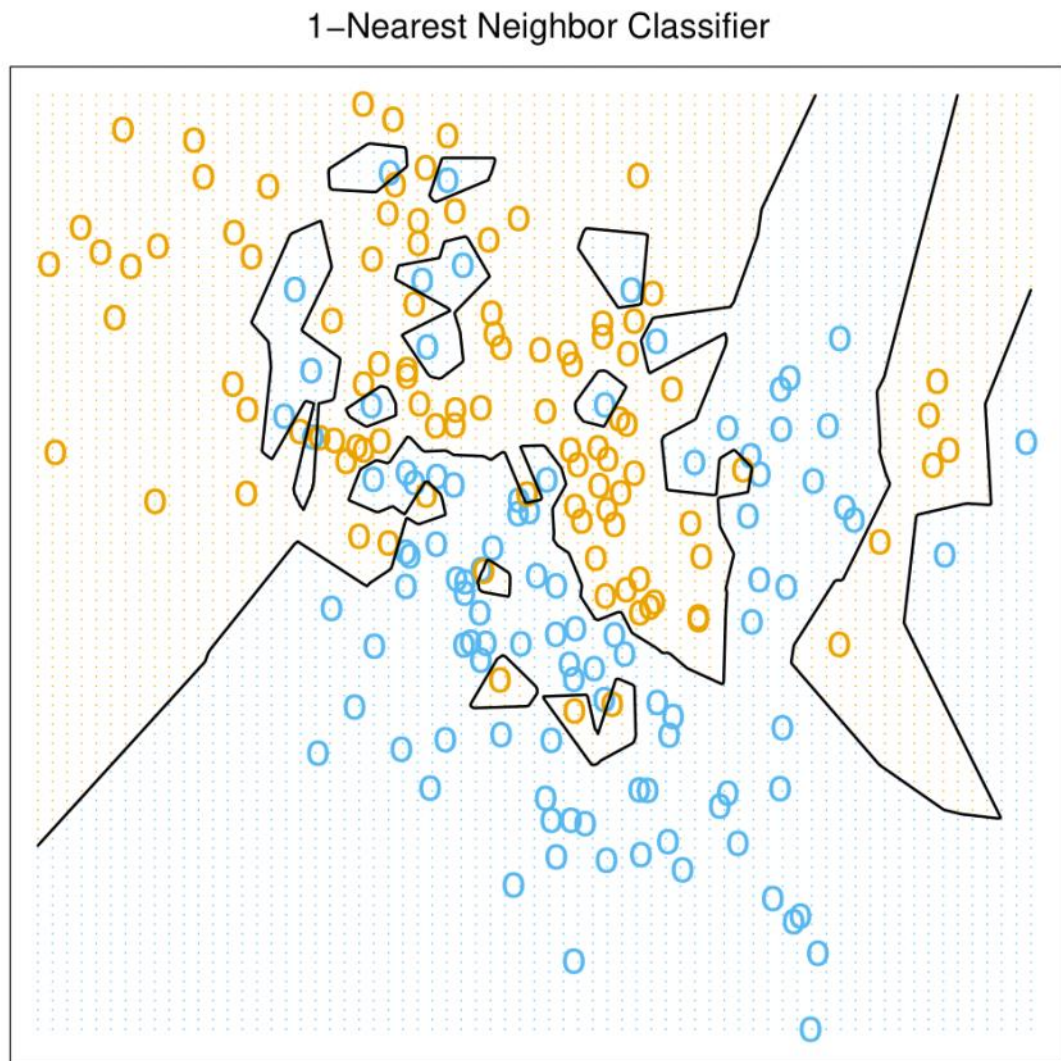


FIGURE 2.3. The same classification example in two dimensions as in Figure 2.1. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1), and then predicted by 1-nearest-neighbor classification.

15-Nearest Neighbor Classifier




FIGURE 2.2. The same classification example in two dimensions as in Figure 2.1. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1) and then fit by 15-nearest-neighbor averaging as in (2.8). The predicted class is hence chosen by majority vote amongst the 15-nearest neighbors.